

Problems in ML & Performance Evaluation

Lương Thái Lê

Outline of the Lecture

- 1. Learning Machine**
- 2. Problems in ML**
 - **Emprical Risk Minimize**
 - **Feature Engineering**
 - **Over fitting**
- 3. ML model Performance Evaluation**
- 4. Evaluation metrics**

Learning Machine

- A learning machine capable of implementing a set of functions

$$f(x, w), w \in \Omega$$

- The learning problem is to choose from the given set of functions the one which best approximates the supervisor's response.
 - The selection is based on training samples $(x_i, y_i), i = 1, \dots, l$
 - => Need to choose and optimize (depend on concrete problem)

Outline of the Lecture

1. Learning Machine

2. Problems in ML


- Empirical Risk Minimize
- Feature Engineering
- Over fitting

3. ML model Performance Evaluation

4. Evaluation metrics

Regression Example: Find road surface

- Given the points, estimate parameters
- Data/Feature
 - dimension ($p=2$) $x_i = (x_{i1}, x_{i2}, \dots x_{ip})^T$
 - # training samples $(x_1, y_1) \dots (x_l, y_l)$
 - parameters $\alpha = (\alpha_0, \alpha_1, \dots \alpha_p)^T$
- Evaluation Metric
 - How good is the fitted plane?



Loss Function

- To choose the best function, it makes sense to **minimize** a loss between the response of the supervisor and the learning machine, given an input x

$$L(y, f(x, \alpha))$$

- Since we want to minimize the loss over *all* samples, we are interested in minimizing the *expected* loss

$$R(\alpha) = \int L(y, f(x, \alpha)) dF(x, y)$$

- $R(\alpha)$ is called Risk function
- $F(x, y)$ is the joint probability distribution function

=> Find $f(x, \alpha^*)$ that minimize $R(\alpha)$ with the only available information is contained in the training set: (x_i, y_i) , $i = 1, \dots, l$

Empirical Risk Minimization Principle

$$R(\alpha) = \int L(y, f(x, \alpha)) dF(x, y)$$

- The risk functional is replaced by the *empirical risk functional*

$$R_{emp}(\alpha) = \frac{1}{l} \sum_{i=1}^l L(y_i, f(x_i, \alpha))$$

=> find $f(x, \alpha^*)$ that minimize $R(\alpha)$ over class of function $f(x, \alpha)$

Loss function: A Probabilistic View

- $L(y, f(x, \alpha)) = \sum_{i=1}^l (y_i - f(x_i, \alpha)) = \sum_{i=1}^l (y_i - \alpha_0 - \sum_{j=1}^p x_{ij} \alpha_j)$
- Let the noise $||\varepsilon|| = ||L(y, f(x, \alpha))||$
- If we model the noise as zero mean Gaussian random variable with variance δ^2 , the distribution is:

$$p(\varepsilon_i) = \frac{1}{\sqrt{2\pi}\delta} \exp\left(-\frac{\varepsilon_i^2}{2\delta^2}\right) = \frac{1}{\sqrt{2\pi}\delta} \exp\left(-\frac{(y_i - \alpha_0 - \sum_{j=1}^p x_{ij} \alpha_j)^2}{2\delta^2}\right)$$
$$= \frac{1}{\sqrt{2\pi}\delta} \exp\left(-\frac{(y_i - \alpha^T x_i)^2}{2\delta^2}\right)$$

$$\Rightarrow p(\varepsilon) = \prod_1^l p(\varepsilon_i) = \frac{1}{(\sqrt{2\pi}\delta)^l} \exp - \frac{1}{2} \left(\frac{\sum_{i=1}^l (y_i - \alpha^T x_i)^2}{\delta^2} \right)$$

Likelihood function

$$\Rightarrow p(\varepsilon) = p(\varepsilon_1, \varepsilon_2, \dots, \varepsilon_l) = \frac{1}{(\sqrt{2\pi}\delta)^l} \exp - \frac{1}{2} \left(\frac{\|y - \alpha^T x\|^2}{\delta^2} \right)$$

We can view this joint probability as a function of the parameters

$$L(\alpha) = p(\varepsilon_1, \varepsilon_2, \dots, \varepsilon_l | \alpha) = \frac{1}{(\sqrt{2\pi}\delta)^l} \exp - \frac{1}{2} \left(\frac{\|y - \alpha^T x\|}{\delta} \right)^2$$

=> Need to maximum Likelihood function

Maximum Likelihood Function

- Maximize the likelihood over all available samples

$$\alpha^* = \operatorname{argmax}_{\alpha} p(\varepsilon_1, \varepsilon_2, \dots, \varepsilon_l | \alpha) = \operatorname{argmax}_{\alpha} \frac{1}{(\sqrt{2\pi}\delta)^l} \exp -\frac{1}{2} \left(\frac{\|y - \alpha^T x\|}{\delta} \right)^2$$

- Since log is a monotonic function, often log-likelihood is used:

$$\alpha^* = \operatorname{argmax}_{\alpha} \left(-\sum_{i=1}^l (y_i - \alpha^T x_i)^2 + \text{const} \right)$$

=> That is where Gradient Descent comes in

Outline of the Lecture

1. Learning Machine

2. Problems in ML

- Empirical Risk Minimize
- **Feature Engineering**
- **Over fitting**

3. **ML model Performance Evaluation**

4. **Evaluation metrics**

Feature Engineering

- Features are individual independent variables that act as an input in your system. In simpler feature is a column of data in your input dataset. Ex: Age, Sex, Income...
- Two types of feature:
 - Categorical: has little values, such as: Sex, Class of ticket (Economy, Premium...), Color,...
 - Numerical: has continuous/discrete values: Age, Price, Name...
- Can create a new feature from a root feature to improve learning
 - Name: Mr John May => create new feature: Tittle (Mr, Miss,...)
- Group features with few values into a generic attribute
 - Tittle with few values, Ex: Rev, Dr, Capital... => Others

Outline of the Lecture

1. Learning Machine


2. Problems in ML

- Empirical Risk Minimize
- Feature Engineering
- **Over fitting**

3. **ML model Performance Evaluation**

4. **Evaluation metrics**

Over fitting



- **Definition**

An objective function that be learnt F will be said to overfit a learning dataset if there exists another objective function F' such that:


- *F' is less suitable (and less accurate) than F for the training set, but*
- *F' is more accurate than F for the entire dataset (including examples used in future)*

- **Reasons of Over-fitting:**

- Error (noise) in the training set (due to the collection/construction data process)
- The number of learning examples is too small to represent the entire examples set of the problems

=> Preferably choose the simplest objective function that fits (nonnecessarily perfect) with training examples

An Over-fitting Example



[*Mitchell, 1997*]

Outline of the Lecture

1. Learning Machine

2. Problems in ML

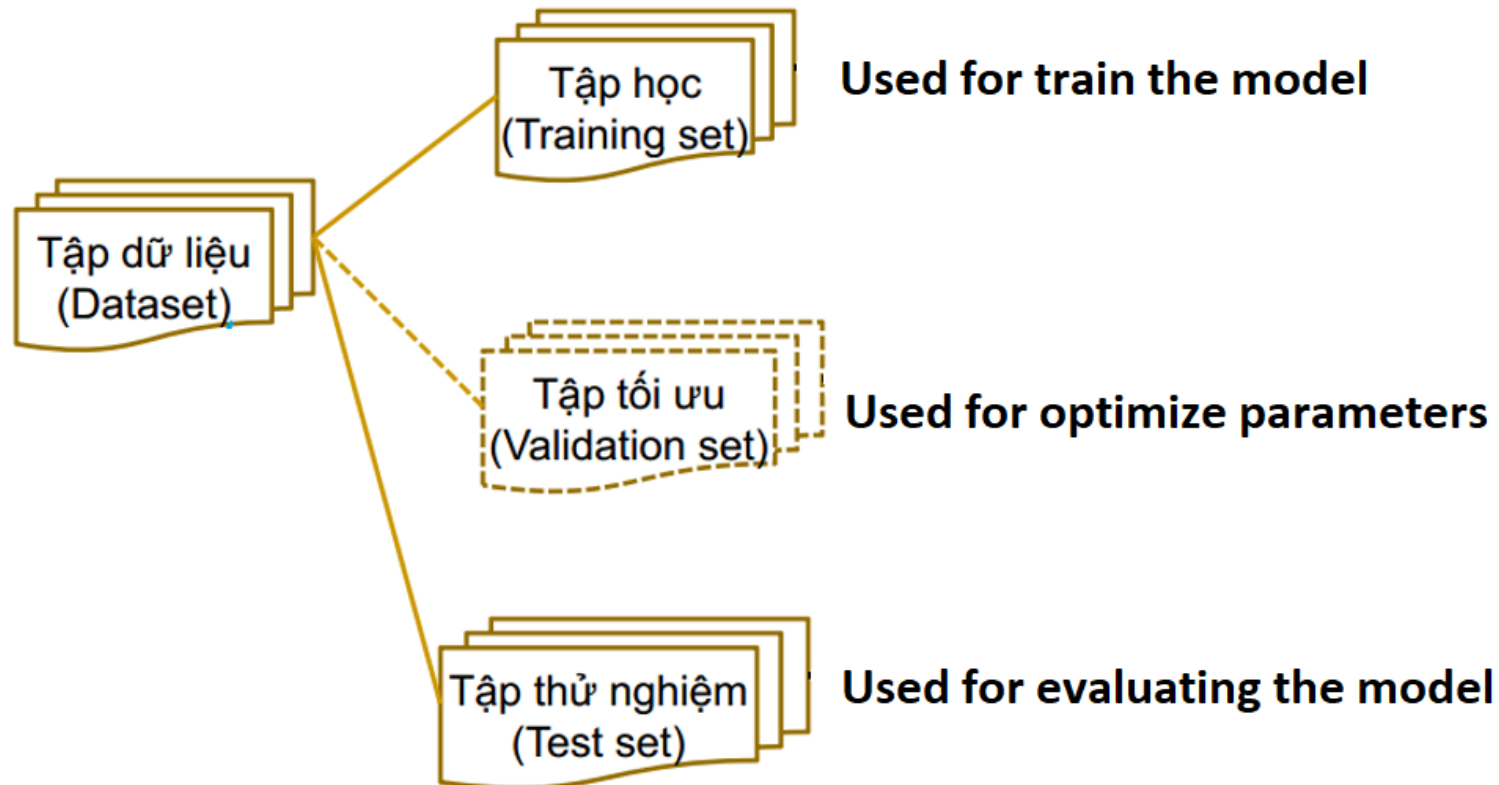
- Empirical Risk Minimize
- Feature Engineering
- Over fitting

3. ML model Performance Evaluation

4. Evaluation metrics

The Model Performance evaluation (1)

- Evaluation of machine learning system performance is often performed empirically, rather than analytically.



The Model Performance evaluation (2)

- The performance of system depends not only on the machine learning algorithms are used, but also depends on:
 - Class distribution
 - Cost of misclassification
 - Size of the training set
 - Size of the test set
- How to obtain a reliable assessment of system performance?
 - The larger the training set, the better the performance of the learning system
 - The larger the test set, the more accurate the evaluation
 - Problem: It is very difficult (rarely) to obtain (very) large data sets

Evaluation Methods

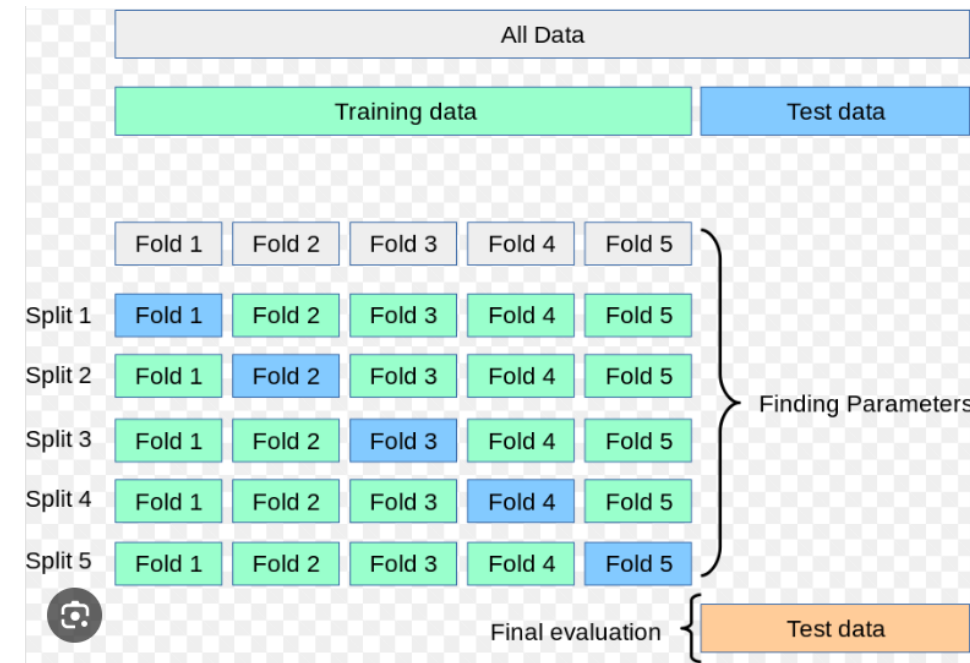
- Hold-out
- Stratified sampling
- Repeated hold-out
- Cross-validation
 - k -fold
 - Leave-one-out
- Bootstrap sampling

Hold-out

- Data set is divided into 2 parts:
 - Training set: D_{train}
 - Test set: D_{test}
- Requirements:
 - Any example in D_{test} is not used in training process
 - Any example in D_{train} is not used in the model evaluation process
- Popular: $|D_{\text{train}}| = 2/3 \cdot |D|$; $|D_{\text{test}}| = 1/3 \cdot |D|$
- Suitable for large set of examples D

Cross Validation – k fold

- The entire set of examples D is divided into k non-intersecting subsets (referred to as “fold”) of approximately the same size
- Each time (of k) iterations, a subset is used as the D_{test} , and $(k-1)$ the remaining subset is used as the D_{train} .
- k error values (each corresponding to a fold) are averaged to get the overall error value
- Popular: $k=10$; or $k=5$



Bootstrap Sampling

- Bootstrap sampling method uses repeated sampling to create a training set
 - Suppose the whole set D consists of n examples
 - From the set D , randomly select an example x (but do not remove x from D)
 - Put example x in the training set: $D_{train} = D_{train} \cup x$
 - Repeat the above 2 steps n times
- Use D_{train} to train the model
- $D_{test} = \{z \in D; z \notin D_{train}\}$ *used for test the model*

Evaluation Criteria

- Accuracy
 - Predictability (classification) of the (trained) model with respect to test instances
- Efficiency
 - Cost of time and resources (memory) required for model training and testing
- Robustness
 - The system's ability to handle (tolerable) noise (error) or missing value
- Scalability
 - How does system performance (e.g. learning/classification rate) change with respect to the size of the data set?
- Interpretability
 - Understanding (for the user) of the system's results and operations easily
- Complexity
 - The complexity of the system model (objective function) learned

Outline of the Lecture

1. Learning Machine
2. Problems in ML
 - Empirical Risk Minimize
 - Feature Engineering
 - Over fitting
3. ML model Performance Evaluation
- 4. Evaluation metrics**

Accuracy

- Show the accuracy of the model when solving the problem
- For the classification problem:

$$Accuracy = \frac{1}{|D_{test}|} \sum_{x \in D_{test}} id(m(x)r(x)) \quad id(a, b) = \begin{cases} 1 & \text{if } a = b \\ 0 & \text{otherwise} \end{cases}$$

- $m(x)$ is the class that the model predict for example x
- $r(x)$ is the real class
- For the regression problem:

$$Error = \frac{1}{|D_{test}|} \sum_{x \in D_{test}} |m(x) - r(x)|$$

- $m(x)$ is the prediction of the model for x
- $r(x)$ is the real output of x

Confuse Matrix (contingency table)

- Only use for classification problem
- **TP**: Number of examples belonging to class c correctly classified into class c
- **TN**: Number of examples that do not belong to class c that be determined exactly
- **FP**: Number of examples that are not in class c isclassified in class c
- **FN**: Number of examples belong to class c but be classified in other class

For class c		Predicted	
		Negative (N) -	Positive (P) +
Actual	Negative -	True Negatives (TN)	False Positives (FP) Type I error
	Positive +	False Negatives (FN) Type II error	True Positives (TP)

Precision and Recall for each class c

- Often used in text classification
- Precision for class c_i :
The total number of examples in class c_i correctly classified divided by the total number of examples classified in class c_i by the model
- Recall for class c_i :
The total number of examples correctly classified by the model into class c_i divided by the total number of examples actually in class c_i

$$Precision(c_i) = \frac{TP_i}{TP_i + FP_i}$$

$$Recall(c_i) = \frac{TP_i}{TP_i + FN_i}$$

Precision and Recall for over all classes

- Assume that the model classifies data into a set of classes $\mathcal{C} = \{c_i\}_{i=1}^n$ and, we get TP_i , TN_i , FP_i , FN_i for each c_i
- Micro - averaging:

$$Precision = \frac{\sum_{i=1}^{|\mathcal{C}|} TP_i}{\sum_{i=1}^{|\mathcal{C}|} (TP_i + FP_i)}$$

$$Recall = \frac{\sum_{i=1}^{|\mathcal{C}|} TP_i}{\sum_{i=1}^{|\mathcal{C}|} (TP_i + FN_i)}$$

- Macro - averaging:

$$Precision = \frac{\sum_{i=1}^{|\mathcal{C}|} Precision(c_i)}{|\mathcal{C}|}$$

$$Recall = \frac{\sum_{i=1}^{|\mathcal{C}|} Recall(c_i)}{|\mathcal{C}|}$$

F_1

- F_1 is a harmonic mean of Precision and Recall

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}}$$

- F_1 tends to take the closest value, whichever is the smaller of the two Precision and Recall values
- F_1 có giá trị lớn nếu cả 2 giá trị Precision và Recall đều lớn

Q&A - Thank you!