

Decision Trees

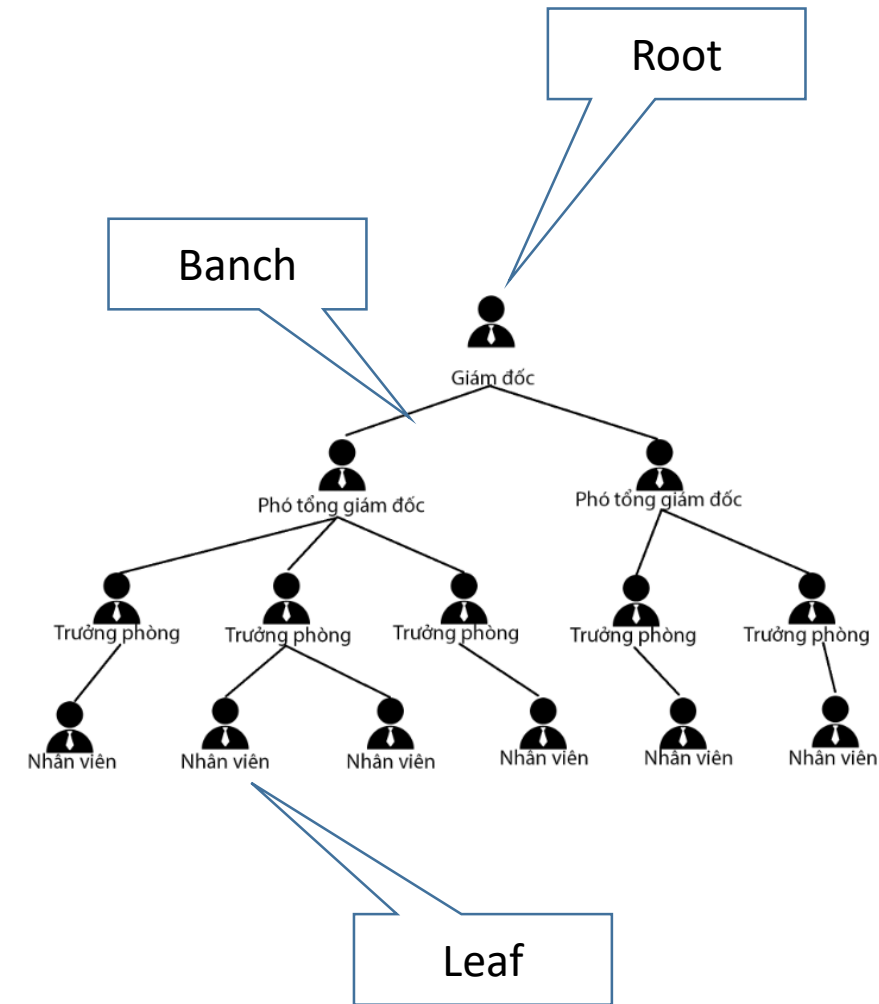
Lương Thái Lê

Outline of the Lecture

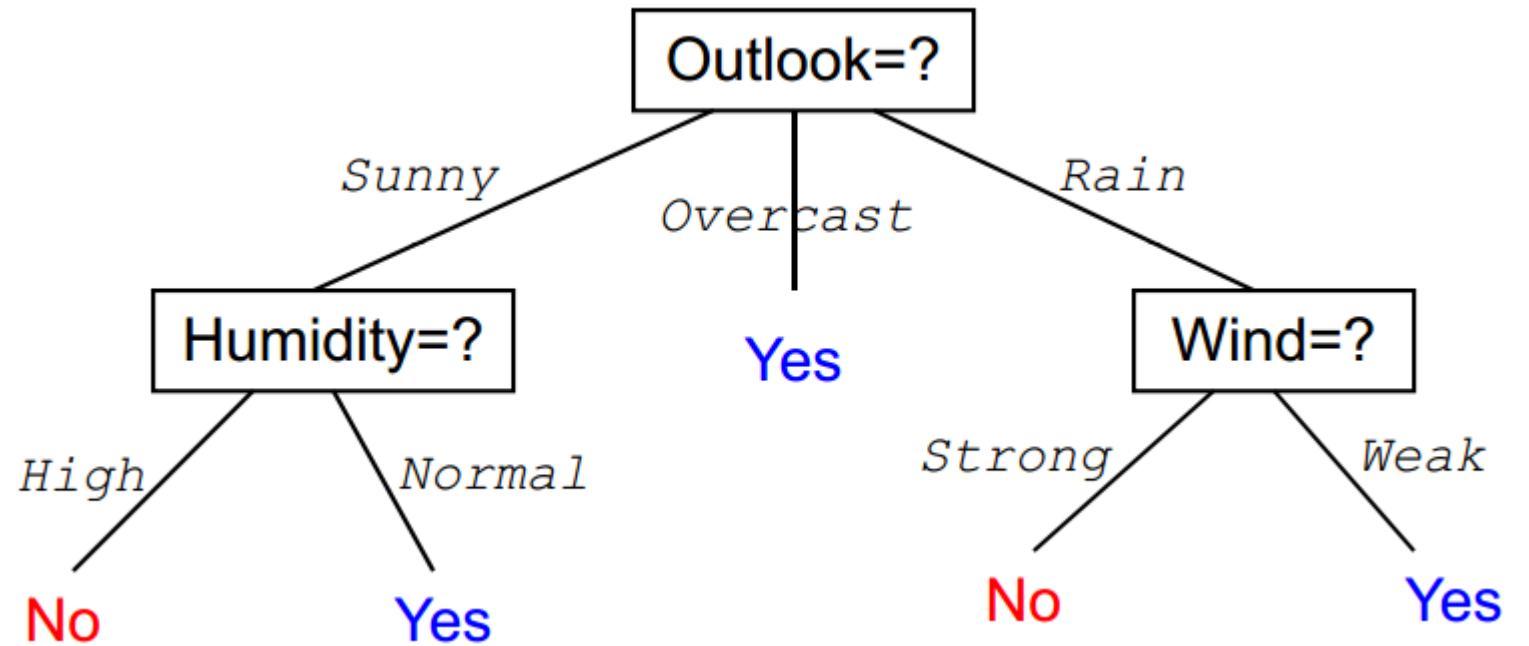
- 1. Introduction of Decision Trees (DT)**
- 2. DT Algorithms**
- 3. Choose the Best Features**
 - Information Gain**
 - Example**

Decision Tree (DT) Introduction

- DT is a supervised learning method – classification
- DT learns a classification function represented by a decision tree
- Can be presented by a set of rules IF – THEN
- Can perform even with noise data
- As one of the most common inductive learning methods
- Successfully applied in many application problems
 - Ex: Spam email filtering...



A DT: Example



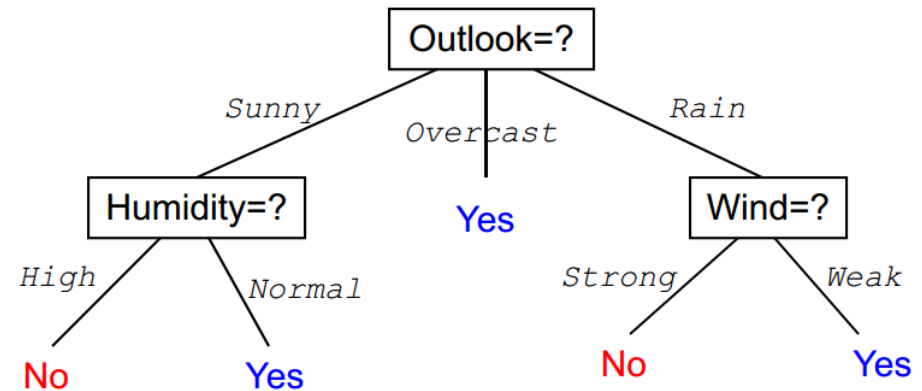
- (Outlook=Overcast, Temperature=Hot, Humidity=High, Wind=Weak) → Yes
- (Outlook=Rain, Temperature=Mild, Humidity=High, Wind=Strong) → No
- (Outlook=Sunny, Temperature=Hot, Humidity=High, Wind=Strong) → No

Represent a DT (1)

- Each internal node represents an attribute to be tested for the examples.
- Each branch from a node corresponds to a possible value of the attribute associated with that node
- Each leaf node represents one class c_i in the set of class C
- A learned DT will classify for an example, by traversing the tree from the root node to a leaf node
 - => The class label associated with that leaf node will be assigned to the example to be classified

Represent a DT (2)

- A DT represents a disjunction of combinations of constraints for the attribute values of the examples
 - Each path from the root node to a leaf node corresponds to a combination of attribute tests



$[(\text{Outlook}=\text{Sunny}) \wedge (\text{Humidity}=\text{Normal})] \vee$
 $(\text{Outlook}=\text{Overcast}) \vee$
 $[(\text{Outlook}=\text{Rain}) \wedge (\text{Wind}=\text{Weak})]$

DT – Problem Setting


- Set of possible instances X :
 - each instance x in X is a feature vector
 - $x = \langle x_1, x_2, \dots, x_n \rangle$; Ex: $\langle \text{Humidity}=\text{low}, \text{Win}=\text{weak}, \text{Outlook}=\text{rain}, \text{Temp}=\text{hot} \rangle$
- Unknown target function: $f: X \rightarrow Y$
 - $y \in Y$; $y = 1$ if we play tennis on this day, else $y = 0$
- Set of function hypotheses $H = \{h \mid h: X \rightarrow Y\}$
 - each hypothesis h is a decision tree

- **Input:**

- Training Examples: $\{ \langle x^{(i)}, y^{(i)} \rangle \}$ of unknown target function f

- **Output:**

- Hypothesis $h \in H$ that best approximates f



Top – down Induction of Decision Trees

[ID3, C4.5, Quinlan]

node = Root

Main loop:

- 1. $A \leftarrow$ the best decision attribute (feature) for next node*
- 2. Assign A as decision attribute for node*
- 3. For each value of A , create descendant of node*
- 4. Sort training examples to leaf nodes*
- 5. If training examples perfectly classified, then STOP else iterate over new leaf node*

Which feature (attribute) is the best?

ID3 Pseudocode (Quinlan - 1979)

ID3 alg (Training_Set, Class_Labels, Attributes)

{

Create the Root node of the decision tree

If all examples of Training_Set belong to the same class c , Return Decision tree with a Root node is labeled c

*If the set Attributes is empty, Return Decision Tree with a Root node attached to a class label \equiv **Majority_Class_Label**(Training_Set)*

*$A \leftarrow$ The attribute in the Attributes set has the "**best**" classifier for Training_Set*

Test Attribute for Root node $\leftarrow A$

For each possible values v of the attribute A

Add a new branch under the Root node, corresponding to the case: "The value of A is v "

Determine $Training_Set_v = \{Instance\ x | x \subseteq Training_Set, x_A = v\}$

If ($Training_Set_v = \emptyset$) then

*Create a leaf node with class label = **Majority_Class_Label**(Training_Set)*

Attach this leaf node to the newly created branch

Else Append to the new created branch a subtree generated by $ID3_alg(Training_Set_v, Class_Labels, \{Attributes\} \setminus \{A\})$

Return Root

}

Choose the Best Attribute


- How to evaluate an attribute's ability to separate learning examples by their class label?

⇒ Use a statistical evaluation

Information Gain

- Example:

Which Attribute will be chosen, A_1 or A_2 ?



PlayTennis: training examples

| Day | Outlook | Temperature | Humidity | Wind | PlayTennis |
|-----|----------|-------------|----------|--------|------------|
| D1 | Sunny | Hot | High | Weak | No |
| D2 | Sunny | Hot | High | Strong | No |
| D3 | Overcast | Hot | High | Weak | Yes |
| D4 | Rain | Mild | High | Weak | Yes |
| D5 | Rain | Cool | Normal | Weak | Yes |
| D6 | Rain | Cool | Normal | Strong | No |
| D7 | Overcast | Cool | Normal | Strong | Yes |
| D8 | Sunny | Mild | High | Weak | No |
| D9 | Sunny | Cool | Normal | Weak | Yes |
| D10 | Rain | Mild | Normal | Weak | Yes |
| D11 | Sunny | Mild | Normal | Strong | Yes |
| D12 | Overcast | Mild | High | Strong | Yes |
| D13 | Overcast | Hot | Normal | Weak | Yes |
| D14 | Rain | Mild | High | Strong | No |

Entropy

- To evaluate the heterogeneity/impurity of a set
- Entropy of the set S for classification with k classes

$$Entropy(S) = \sum_{i=1}^k -p_i \log_2 p_i$$

where p_i is the proportion of examples in the set S that belong to class i , and $0 \cdot \log_2 0 = 0$

- Entropy of the set S for classification with 2 classes

$$H(S) \equiv -(p_1 \log_2 p_1) - (p_2 \log_2 p_2)$$

- The meaning of entropy in the field of Information Theory
 - The entropy of the set S indicates the number of bits required to encode the class of an element randomly drawn from the set S .

Entropy – Example with 2 classes

- S includes 14 examples, of which 9 belong to class c_1 (Yes) and 5 examples belong to class c_2 (No)

$$\Rightarrow \text{Entropy}(S) = -\left(\frac{9}{14}\right) \cdot \log_2\left(\frac{9}{14}\right) - \left(\frac{5}{14}\right) \cdot \log_2\left(\frac{5}{14}\right) \approx 0,94$$

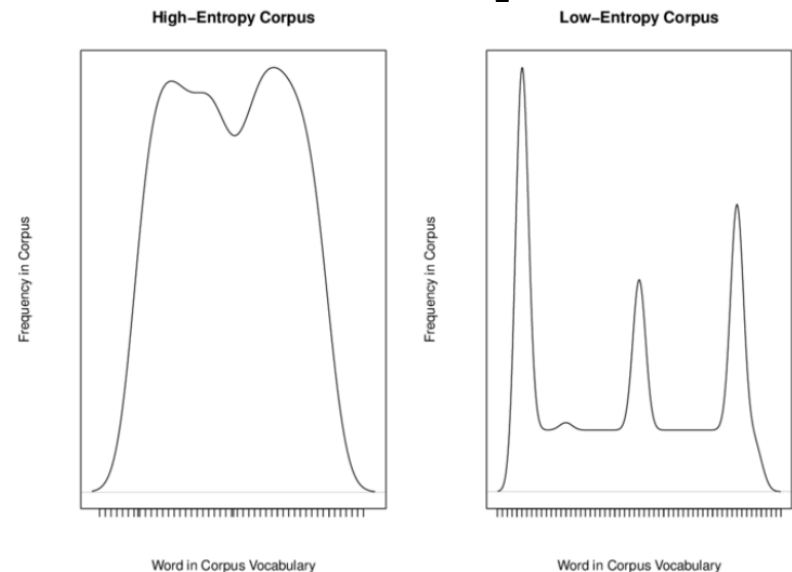

- Entropy = 0, if all examples belong to the same class (c_1 or c_2)
- Entropy = 1, if the number of examples belonging to class c_1 is equal to the number of examples belonging to class c_2
- Entropy = a value in the range (0,1), if the number of examples belonging to class c_1 is different from the number of examples belonging to class c_2

\Rightarrow High Entropy:

- x is from a uniform like distribution
- values sampled from it are less predictable

\Rightarrow Low Entropy:

- x is from a varied (peaks and valley) distribution
- values sampled from it are easier predictable



Information Gain

- Information Gain of an attribute for a set of examples:
 - The reduction degree in Entropy by partitioning the examples by the values of that property

- Information Gain of attribute A for the set S:

$$IG(S, A) = Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

where $Values(A)$ is the set of possible values of the attribute A and

$$S_v = \{x | x \in S, x_A = v\}$$

- Meaning of $IG(S, A)$:
 - The number of bits reduced for the class encoding of a random example from the set S, when the value of attribute A is known.

=> The best feature is the feature with highest IG

The learning set S (Mitchell-1998)

| Day | Outlook | Temperature | Humidity | Wind | Play Tennis |
|-----|----------|-------------|----------|--------|-------------|
| D1 | Sunny | Hot | High | Weak | No |
| D2 | Sunny | Hot | High | Strong | No |
| D3 | Overcast | Hot | High | Weak | Yes |
| D4 | Rain | Mild | High | Weak | Yes |
| D5 | Rain | Cool | Normal | Weak | Yes |
| D6 | Rain | Cool | Normal | Strong | No |
| D7 | Overcast | Cool | Normal | Strong | Yes |
| D8 | Sunny | Mild | High | Weak | No |
| D9 | Sunny | Cool | Normal | Weak | Yes |
| D10 | Rain | Mild | Normal | Weak | Yes |
| D11 | Sunny | Mild | Normal | Strong | Yes |
| D12 | Overcast | Mild | High | Strong | Yes |
| D13 | Overcast | Hot | Normal | Weak | Yes |
| D14 | Rain | Mild | High | Strong | No |

Information Example

- Calculate the Information Gain value of the Wind attribute for the learning set S : $IG(S, Wind)$
- The Wind attribute has 2 possible values: Weak and Strong
- $S = \{9 \text{ for Yes, and } 5 \text{ for No}\}$
- $S_{\text{weak}} = \{6 \text{ examples of Yes class and } 2 \text{ examples of No class with value Wind=Weak}\}$
- $S_{\text{strong}} = \{3 \text{ examples of Yes class and } 3 \text{ examples of No class with value Wind=Strong}\}$

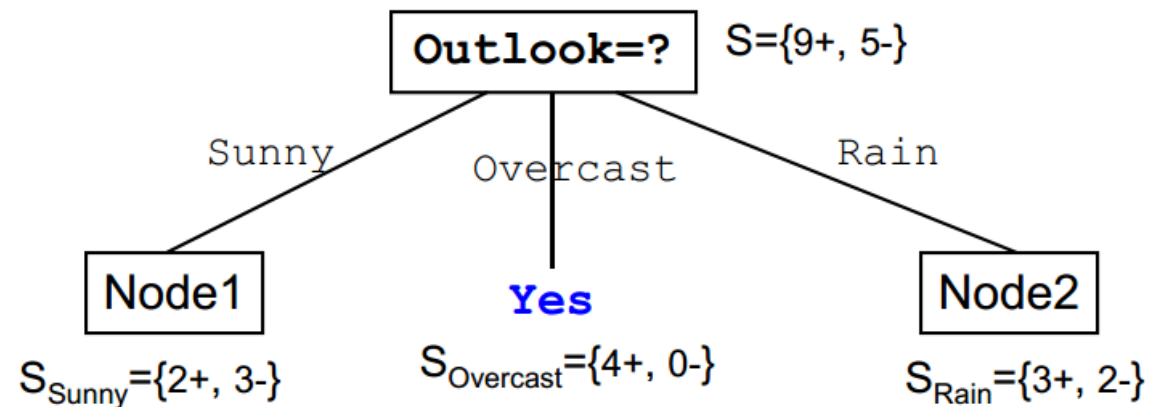
$$\begin{aligned} IG(S, Wind) &= Entropy(S) - \sum_{v \in \{Weak, Strong\}} \frac{|S_v|}{|S|} Entropy(S_v) \\ &= Entropy(S) - (8/14)Entropy(S_{Weak}) - (6/14)Entropy(S_{Strong}) \\ &= 0,94 - (8/14)0,81 - (6/14)1 = 0,048 \end{aligned}$$

Learning a DT – Example (1)

- For the Root, choose the best feature from the set {Outlook, Temperature, Humidity, Wind}
 - $IG(S, \text{Outlook}) = \dots = 0,246$
 - $IG(S, \text{Temperature}) = \dots = 0,029$
 - $IG(S, \text{Humidity}) = \dots = 0,151$
 - $IG(S, \text{Wind}) = \dots = 0,048$

The highest IG

=> Outlook is chosen to be the test feature for the Root



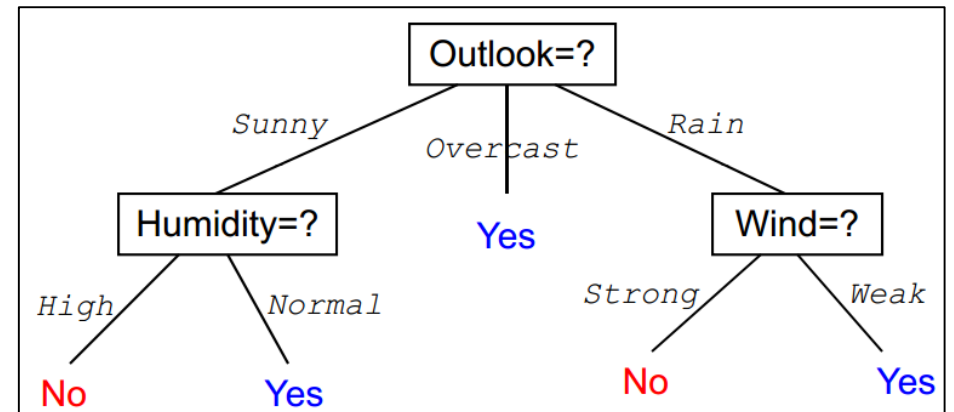
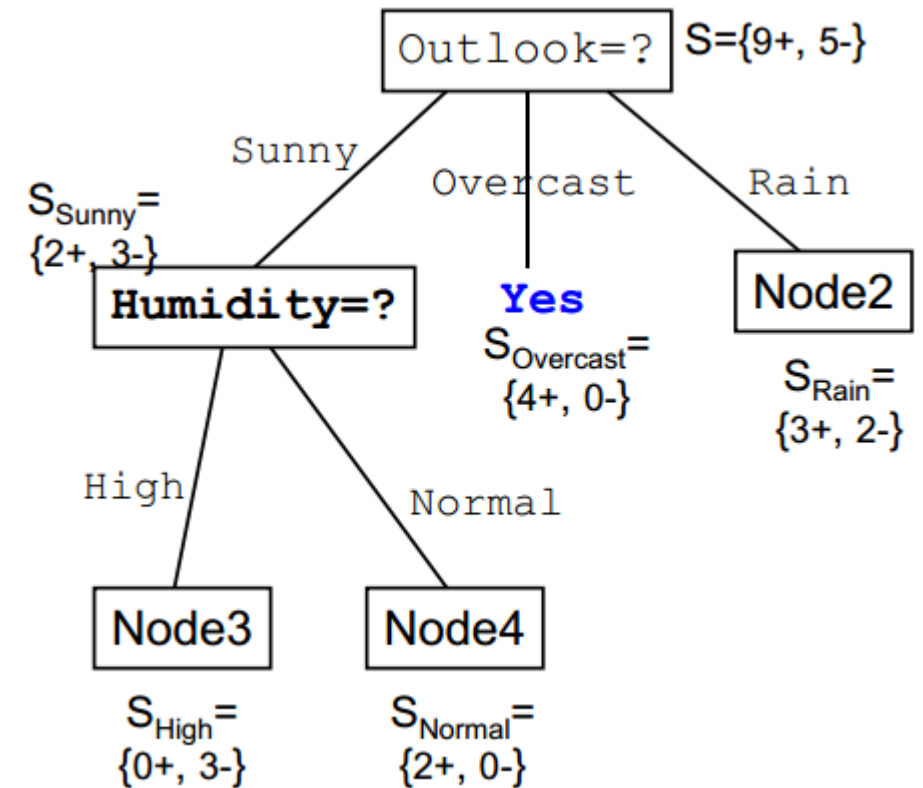
Learning a DT – Example (2)

- For the Node1, choose the best feature from the set {Temperature, Humidity, Wind} to be test feature.

- $IG(S_{Sunny}, Temperature) = \dots = 0,57$
- $IG(S_{Sunny}, Humidity) = \dots = 0,97$
- $IG(S_{Sunny}, Wind) = \dots = 0,57$

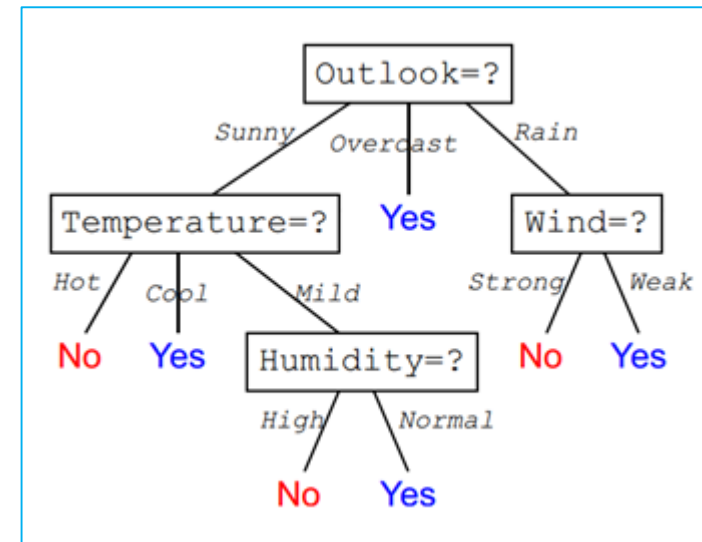
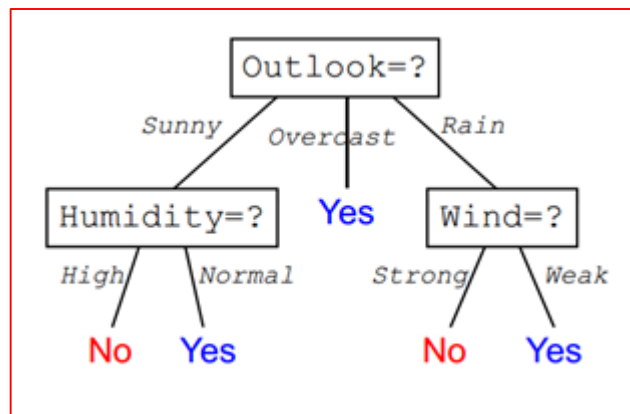
⇒ Choose Humidity for Node1

⇒ Similar, we have Node2, Node3, Node4



Comment on Stragy of ID3

- ID3 searches only one (but not all) decision trees that fit the training examples
 - chooses the first matching decision tree found during its search
- Use Information Gain to choose the best test feature
 - ⇒ bias towards multivalued attributes (Ex: Bank account, ID,...) ⇒ easily to get overfitting
- During the search, ID3 does not perform backtracking
 - ⇒ It is only guaranteed to find a locally optimal solution,



Problems in ID3 that Need to be Solved

- Overfitting
- Handling attributes with continuous value (Age, Price...)
- The more suitable evaluations (better than Information Gain) for determining the test attribute for a node
- Handling missing-value attributes training examples
- Handling attributes with different costs

=> C4.5 can handle all above problems

Overfitting Solving

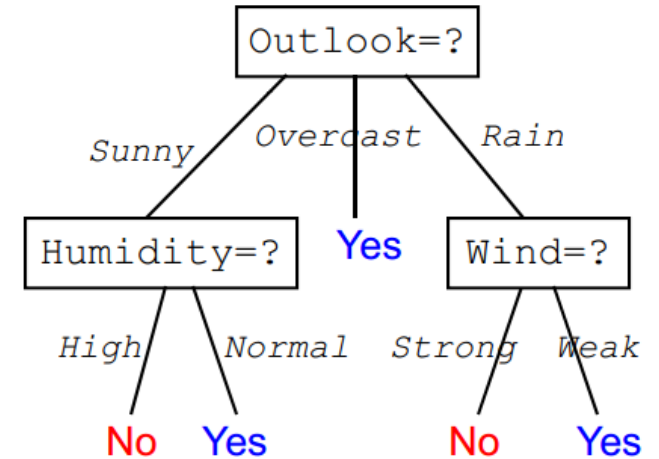
- 2 strategies:
 - Stop learning the decision tree earlier, before it reaches a tree structure that matches perfect classification of the training set
 - ⇒ difficult to decide when to stop
 - Learn the full tree (perfectly suitable for the training set), and then perform the tree pruning process.
 - ⇒ often give better performance in practice
- How to properly prune trees?
 - Evaluation of classifier performance for an validation set
 - Use *reduced-error pruning* and *rule post-pruning*

Reduced-error Pruning

- Each node of the completely tree is checked for pruning
- A node will be pruned if the tree (after pruning that node) achieves no worse performance than the original tree for the validation set.
- Pruning a node includes:
 - Remove all sub-trees associated with pruned node
 - Convert pruned node to a leaf node (classified label)
 - Attached to this leaf node (pruned node) the class label that dominates the training set associated with that node
- Repeat pruning
 - Always select a node that pruning maximizes the likelihood classification of the decision tree for validation set
 - Stop pruning when it reduces the classifiability of decision tree for the validation set

Rule post-pruning

- Convert the complete decision tree learned into a set of corresponding rules
- Reducing each rule (independently of the others) by removing any conditions that do not help bring about an improvement in the classification efficiency of that rule
- Rearrange the reduced rules according to the classifier ability, and use this order for the classification of future examples



IF (Outlook=Sunny) \wedge
 (Humidity=Normal)
THEN (PlayTennis=Yes)

Features with Continuous values

- Need to convert to discrete-valued attributes, by dividing the continuous interval into a set of non-intersecting intervals.
- For the (continuous) attribute A , create a new attribute of binary type A_v such that: A_v is True if $A > v$, and False otherwise.
- How to determine the “best” threshold value v ?
 - Choose the threshold value v that produces the highest Information Gain value
- Example:
 - Sort the learning examples in ascending value for the Temperature
 - Identify learning examples that are contiguous but different from class (Temperature 48 & 60; Temperature 80 & 90)
 $\Rightarrow \text{average}(48,60)=54$; $\text{average}(80,90)=85$
 - There are 2 possible threshold values: Temperature₅₄ and Temperature₈₅
 - The new binary feature Temperature₅₄ is selected, because $IG(S, \text{Temperature}_{54}) > IG(S, \text{Temperature}_{85})$

| | | | | | | |
|-------------|----|----|-----|-----|-----|----|
| Temperature | 40 | 48 | 60 | 72 | 80 | 90 |
| PlayTennis | No | No | Yes | Yes | Yes | No |

Gain Ratio – Another Way choosing the best feature

- → Reduce the effect of attributes with many values

$$SplitInformation(S, A) = - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} \log_2 \frac{|S_v|}{|S|}$$

$$GainRatio(S, A) = \frac{IG(S, A)}{SplitInformation(S, A)}$$

where $Values(A)$ is the set of possible values of the attribute A

$$S_v = \{x | x \in S, x_A = v\}$$

Handling attributes with missing values (1)

- Suppose attribute A is a candidate for the test attribute at node n
- How to deal with the example x has no value for attribute A
- Let S_n be the set of training examples associated with node n that have a value for the attribute A
 - Solution 1: x_A is the most common value for attribute A among the examples belonging to the set S_n
 - Solution 2: x_A is the most common value for attribute A among the examples belonging to the set S_n having the same target class as x

Attributes have different costs

- In some machine learning problems, attributes can be assigned different costs
 - Example: In learning to classify medical diseases, BloodTest has costs \$150, while TemperatureTest costs \$10
 - Tendency to learn cost-based decision trees:
 - Use as many low-cost attributes as possible
 - Only use high-cost attributes when necessary (to help achieve reliable classifications)
- => Using assessments other than IG for test attribute identification

$$\frac{Gain^2(S, A)}{Cost(A)}$$

[Tan and Schlimmer, 1990]

$$\frac{2^{Gain(S, A)} - 1}{(Cost(A) + 1)^w}$$

[Nunez, 1988; 1991]

When to use DT?

- Learning examples are represented by (attribute, value) pairs.
 - Suitable with discrete-valued attributes
 - For attributes with continuous values, it must be discretized
- The objective function whose output is discrete values
 - Example: Classify the examples into the appropriate class
- The training set may contain noise/error
- The training set may contain missing attributes

Q&A - Thank you!